
UCK358E – INTR. TO ARTIFICIAL INTELLIGENCE

SPRING '23

LECTURE 1

INTRODUCTION

Instructor: Asst. Prof. Barış Başpınar

What are all these buzzwords?

Artificial Intelligence

Machine Learning

Big Data

Deep Learning

Neural Networks

Data Science

Pattern Recognition

Data Mining

Here is the dictionary!

Deep Learning = Neural Networks

Neural Networks \subset Machine Learning

Machine Learning \approx Artificial Intelligence

All other buzzwords \approx Machine Learning

Timeline of Ups and Downs

AI: Artificial Intelligence

ML: Machine Learning

- **1950's:** AI/ML is the great future.
- **1970's:** AI became a bad word.
- **1980's:** Neural Networks are the great future.
- **1990's:** Neural Networks are not that good.
- **2010's:** Neural Networks are great after all.
- **Now:** AI/ML is the great future.

Why Learning?



Jeopardy's Watson is a one-task machine. Big task, but one task. (2011)

What the buzzwords have in common

- The same core premise:
 - *Machine Learning*
 - *Artificial Intelligence*
 - *Data Mining*
 - *Pattern Recognition*

“Automated detection of a pattern based on the data ”

Example: Credit Approval

Given the data of an applicant:

age	23 years
gender	male
annual salary	\$30,000
years in residence	1 year
years in job	1 year
current debt	\$15,000

should we extend credit?


When should ML be used?

- ML is the technology of choice when:
 - A **pattern** exists.
 - We cannot pin it down mathematically.
 - We have a representative **data** set.


These criteria led to 3 waves of successful applications

1st Wave: Financial applications

- Market forecasting



- Financial model calibration



- Consumer and corporate credit assessment.

2nd Wave: E-commerce

- Recommender systems (Amazon, fashion, ...)



- Profiling




Famous ML e-commerce problem



US\$1,000,000 Prize for the first 10% improvement (2006-2009)

3rd Wave: Medical Applications

- Medical diagnosis



- Data mining of medical records



The Growth of Machine Learning

- More Data:

Enables us to pin down the pattern better.

gradual → jump

- More Complex Models:

Enables us to capture more complex patterns.

gradual → jump

- More Computation:

Enables us to optimize very complex models.


gradual → jump

ML Success Stories

The last years witnessed a huge surge in the practical impact of ML



speech recognition



object detection



machine translation

What happened?

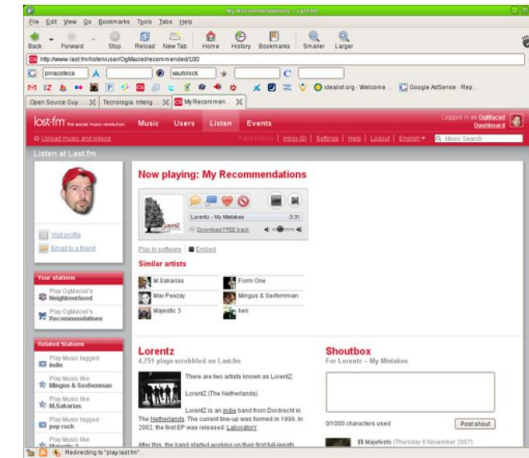
A qualitative change in:

data - models - computation

Different **data**: Total Profiling

Using multiple data sources:

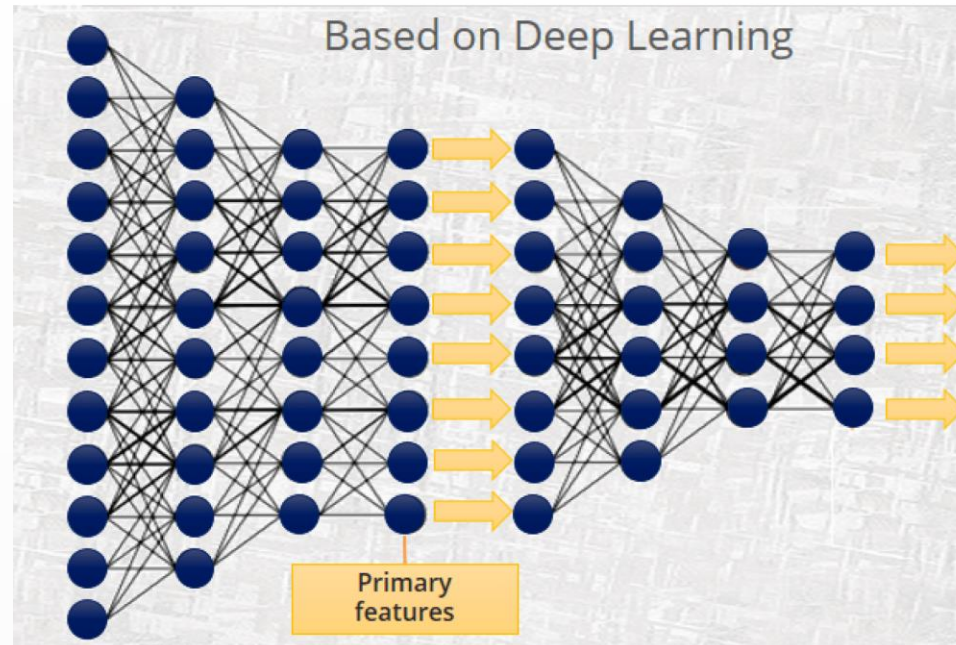
Movie preferences, Facebook posts, Amazon purchases, etc.



to profile a person.

What about the **models**?

A deep Neural Network with millions of parameters



Automated Feature Extraction

Jump in **computation** speed

Commercially available specialized hardware

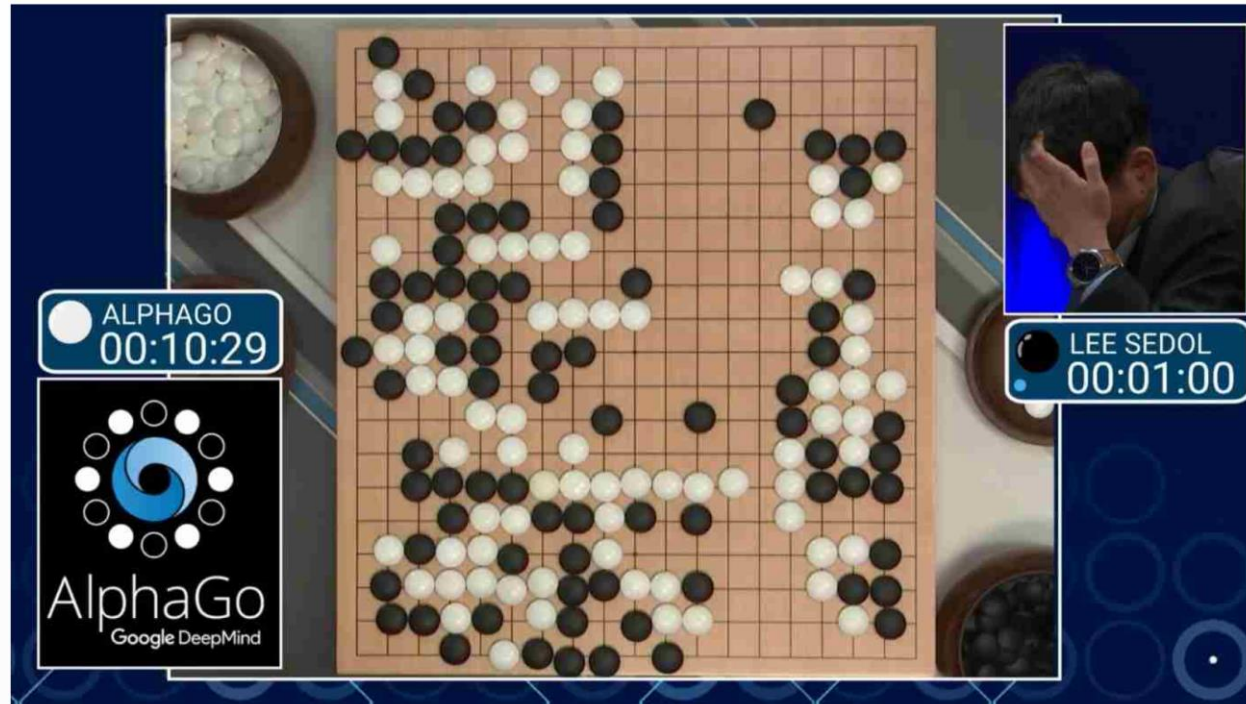


Gain in ML speed is more than 2 orders of magnitude

Real “Intelligence” achieved

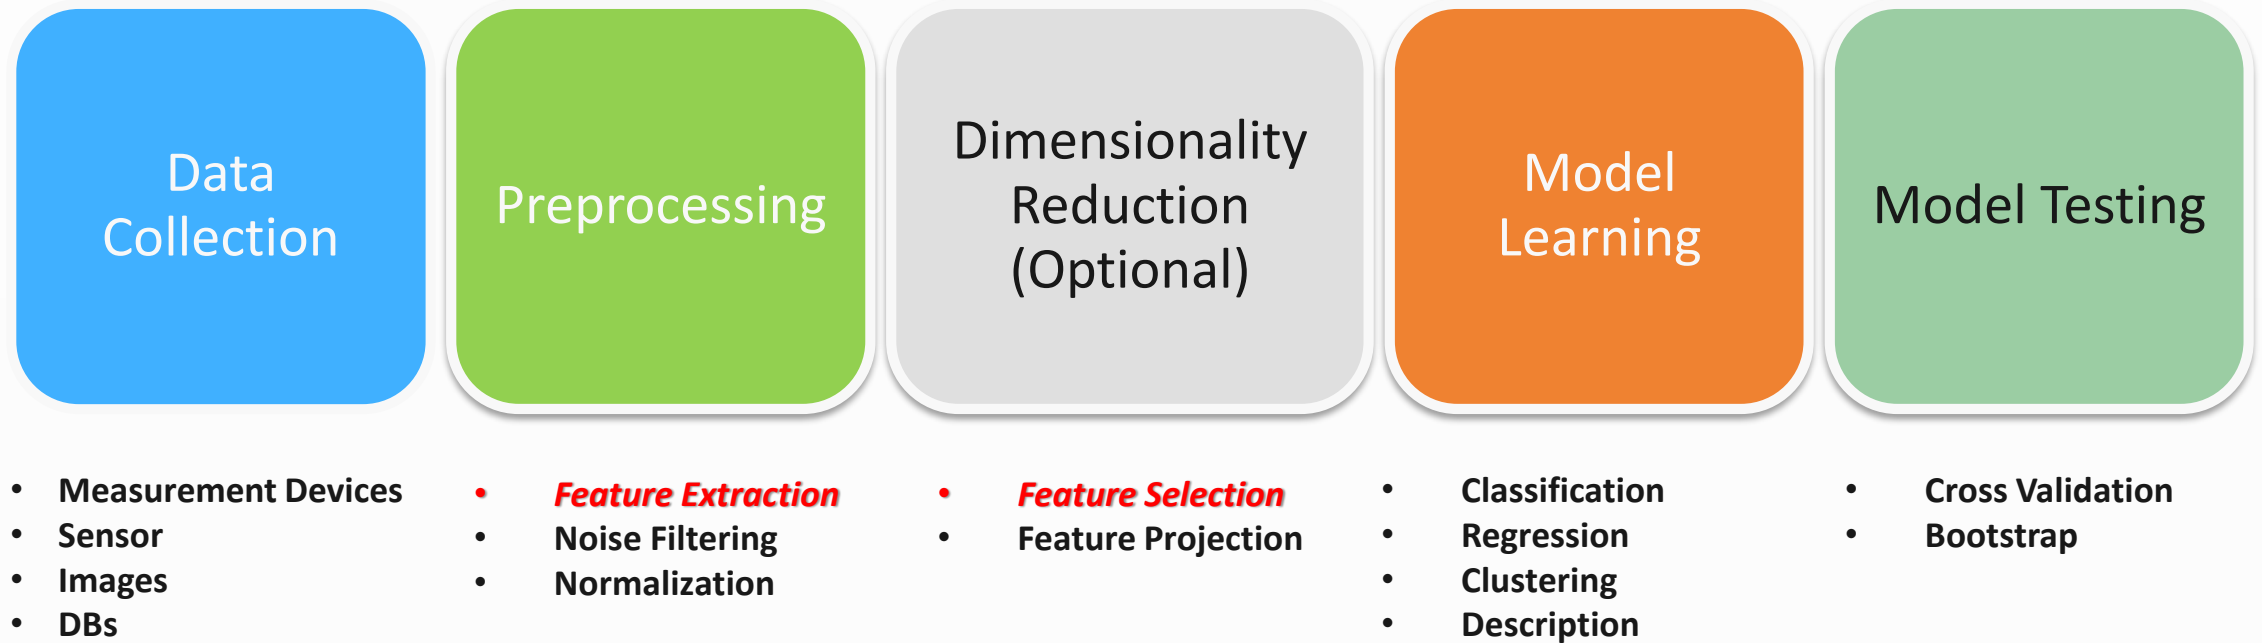
From: Replicating human skills

To: Beating human intelligence




ML system discovered novel strategic moves in the game of **Go** (2015)

Machine Learning Perspective

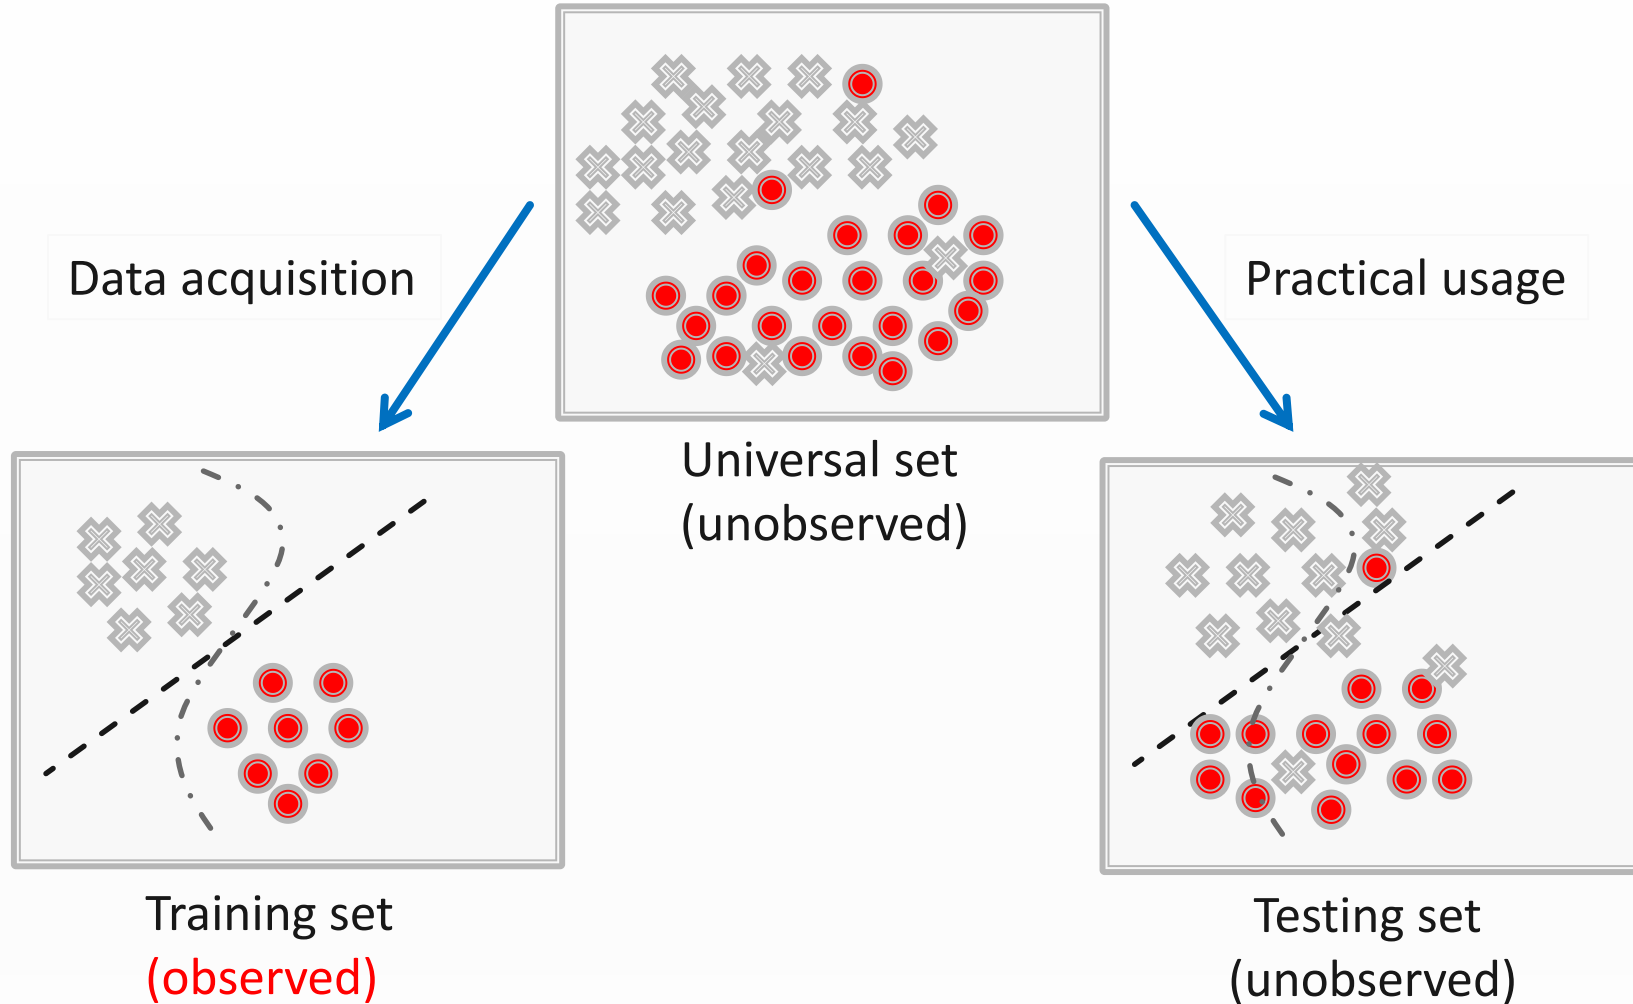


Learning Essentials

- **Machine Learning:** Data -> Train a Model on Data -> Make Predictions on New Data
- **Feature Engineering:** Extracting useful patterns from data that will be used for ML models for classification
- **Feature Learning:** Feature learning algorithms find the common patterns that are important to classify samples and extract them automatically to be used in a classification or regression process.
- **Deep Learning:** New methods and strategies designed to generate deep hierarchies of non-linear features so that Deep architectures with dozens of layers of non-linear hierarchical features can be trained




Training and testing



Training and testing

- Training is the process of making the system able to learn.
- No free lunch rule:
 - Training set and testing set come from the same distribution
 - Need to make some assumptions or bias



Performance

- There are several factors affecting the performance:
 - **Types of training** provided
 - The form and extent of any initial **background knowledge**
 - The **type of feedback** provided
 - The **learning algorithms** used

- Two important factors:
 - Modeling
 - Optimization

Algorithms


- The success of machine learning system also depends on the algorithms.
- The algorithms control the search to find and build the knowledge structures.
- The learning algorithms should extract useful information from training examples.

Algorithms


- **Supervised learning** ($\{x_n \in R^d, y_n \in R\}_{n=1}^N$)
 - Prediction
 - Classification (discrete labels), Regression (real values)
- **Unsupervised learning** ($\{x_n \in R^d\}_{n=1}^N$)
 - Clustering
 - Probability distribution estimation
 - Finding association (in features)
 - Dimension reduction
- **Semi-supervised learning**
- **Reinforcement learning**
 - Decision making (robot, chess machine)

Algorithms


Supervised learning



Unsupervised learning

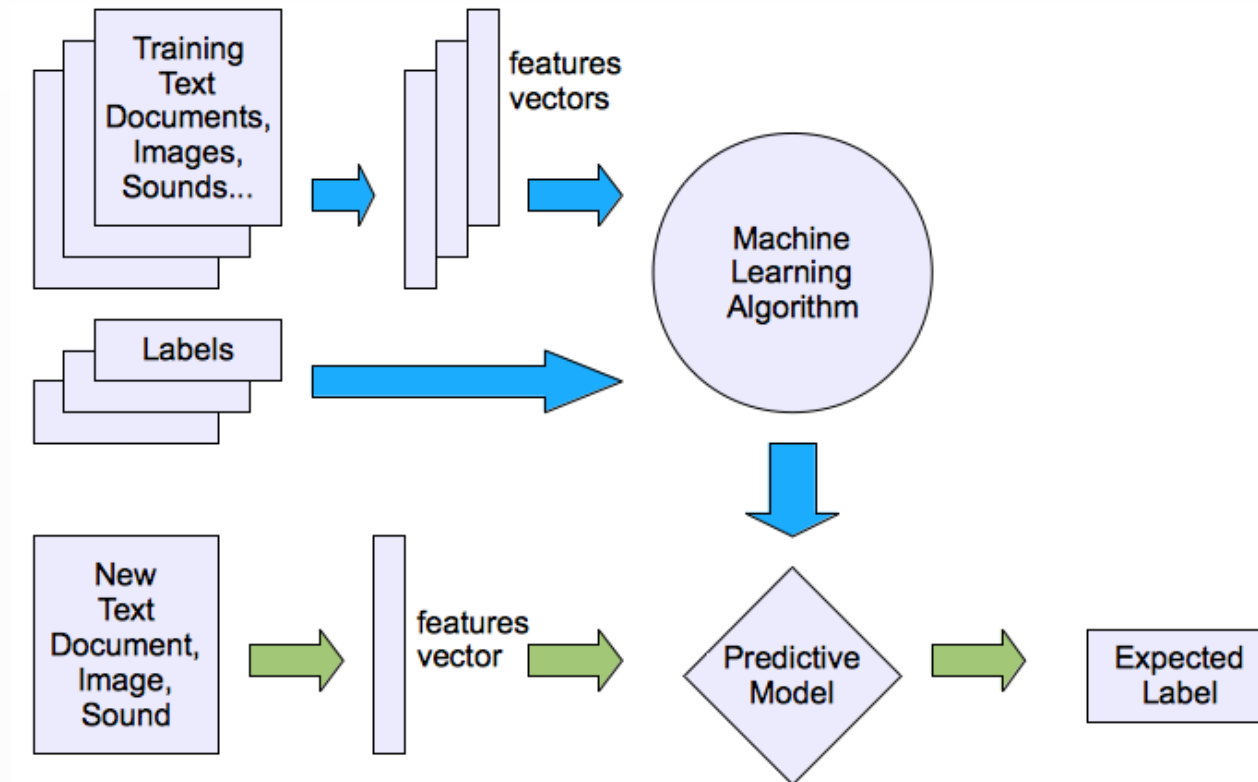


Semi-supervised learning



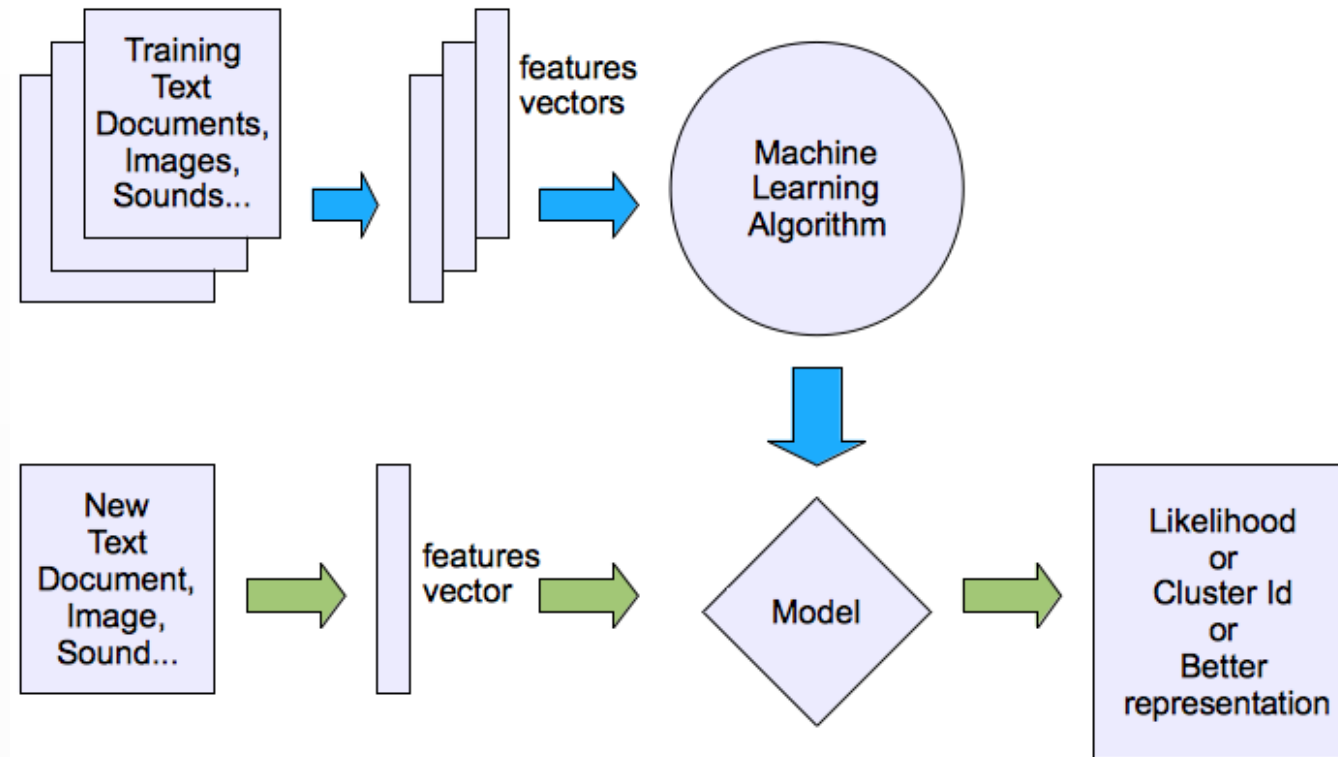
Machine learning structure

- Supervised learning



Machine learning structure

- Unsupervised learning



What are we seeking?

- Supervised: Low E-out or maximize probabilistic terms

$$error = \frac{1}{N} \sum_{n=1}^N [y_n \neq g(x_n)]$$

E-in: for training set


E-out: for testing set

$$E_{out}(g) \leq E_{in}(g) \pm O\left(\sqrt{\frac{d_{VC}}{N} \ln N}\right)$$

- Unsupervised: Minimum quantization error, Minimum distance, MAP, MLE(maximum likelihood estimation)

What are we seeking?

Under-fitting VS. Over-fitting (fixed N)




Learning techniques

- Supervised learning categories and techniques
 - **Linear classifier** (numerical functions)
 - **Parametric** (Probabilistic functions)
 - Naïve Bayes, Gaussian discriminant analysis (GDA), Hidden Markov models (HMM), Probabilistic graphical models
 - **Non-parametric** (Instance-based functions)
 - *K*-nearest neighbors, Kernel regression, Kernel density estimation, Local regression
 - **Non-metric** (Symbolic functions)
 - Classification and regression tree (CART), decision tree
 - **Aggregation**
 - Bagging (bootstrap + aggregation), Adaboost, Random forest

Learning techniques

- Linear classifier




$$g(x_n) = \text{sign}(w^T x_n)$$

, where w is an d -dim vector (learned)

- Techniques:
 - Logistic regression
 - Support vector machine (SVM)
 - Multi-layer perceptron (MLP)

Learning techniques

- Support vector machine (SVM):



$$x_n = [x_{n1}, x_{n2}]$$



$$x_n = [x_{n1}, x_{n2}, x_{n1} * x_{n2}, x_{n1}^2, x_{n2}^2]$$
$$g(x_n) = \text{sign}(w^T x_n)$$



- Non-linear case

- Linear to nonlinear: **Feature transform** and **kernel function**

Learning techniques

- Unsupervised learning categories and techniques
 - **Clustering**
 - K-means clustering
 - Spectral clustering
 - **Density Estimation**
 - Gaussian mixture model (GMM)
 - Graphical models
 - **Dimensionality reduction**
 - Principal component analysis (PCA)
 - Factor analysis

Deep Learning Essentials



```
class Neuron(object):
```

```
# ...
```




```
def forward(inputs):
```

```
    """ assume inputs and weights are 1-D numpy arrays and bias is a number """
```

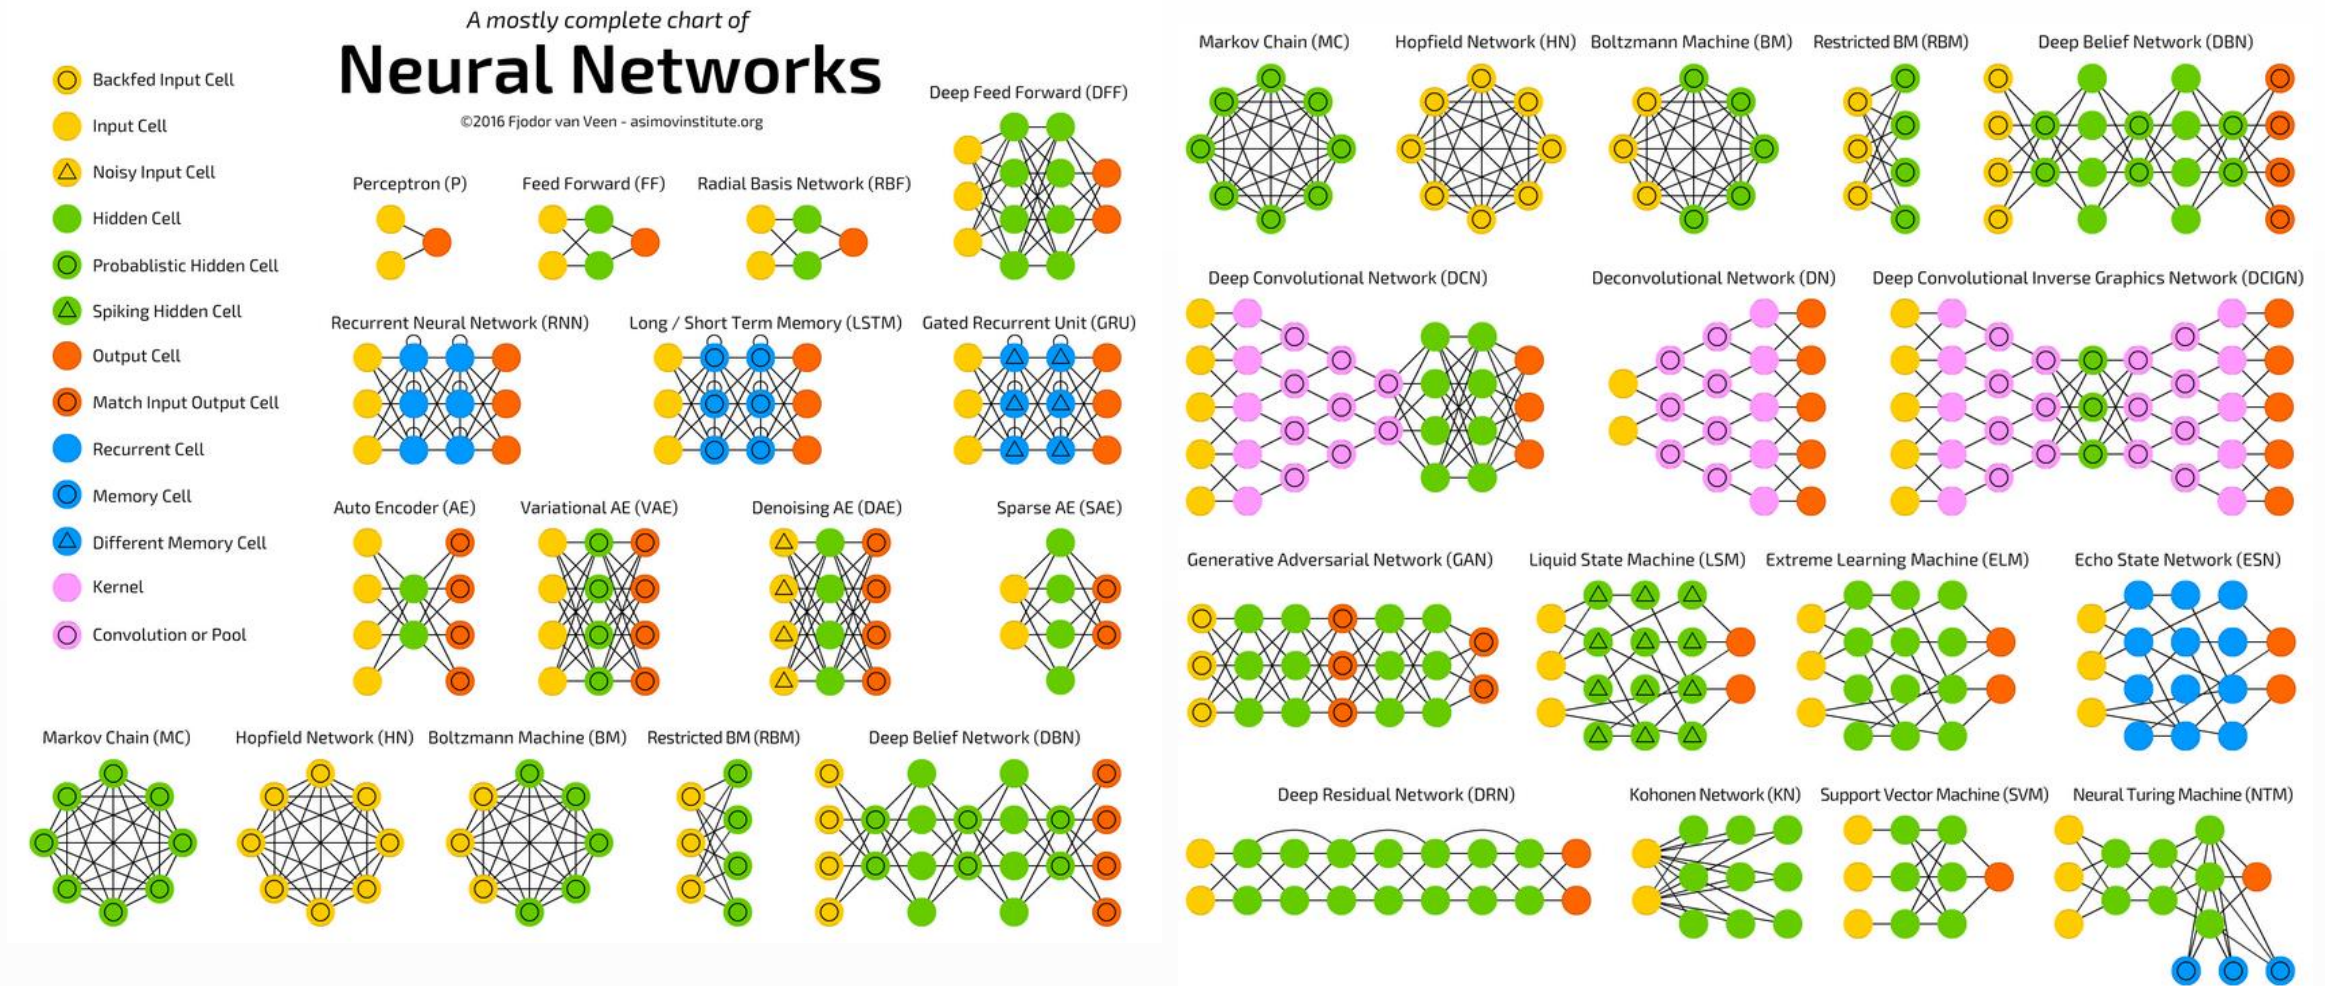
```
    cell_body_sum = np.sum(inputs * self.weights) + self.bias
```

```
    firing_rate = 1.0 / (1.0 + math.exp(-cell_body_sum)) # sigmoid activation function
```

```
    return firing_rate
```



Neural Network Architectures



Applications

- Face detection
- Object detection and recognition
- Image segmentation
- Multimedia event detection
- Economical and commercial usage

References

- W. L. Chao, J. J. Ding, “Integrated Machine Learning Algorithms for Human Age Estimation”, NTU, 2011.
- Y. S. Abu-Mostafa, “Artificial Intelligence: Evolution & Revolution (& Hype!)”, 2018.